

THERMO FISHER SCIENTIFIC - R-01328445

Architecture should make trusted scale easier.

Dear Hiring Team,

Thermo Fisher is not hiring this role merely to choose models or draw cloud diagrams. The harder mandate is to turn distributed AI and data demand into a reusable enterprise capability that remains scalable, secure, observable, cost-aware, and credible in scientific work.

My career has repeatedly placed me at the boundary between emerging technology and operating reality. At Vape-Jet, I lead AI-first transformation across robotics, machine vision, ERP, CRM, quality, support, engineering issue flow, and telemetry. Through DudeWorth, I design AI-readiness and agentic workflow patterns that make data authority, human judgment, evaluation, governance, and adoption explicit. At Compunetics, I led high-reliability engineering and operations systems where controls and traceability were non-negotiable. Amazon strengthened the operating discipline required to make standards work under scale.

I do not claim to have owned Thermo Fisher's exact AWS/Azure, Databricks, Dataiku, vector-database, or UiPath estate. I would earn architectural authority by reading the existing system, partnering with platform and domain experts, and proving fit through one bounded, decision-critical use case before broader standardization.

The candidate vision at russelldudek.github.io/thermo/ makes that reasoning visible through a Scientific Capability Lattice, a detailed 120-day entry plan, and a Load-Bearing Architecture Review. The thesis: AI architecture should be modular at the edge, governed at the core, and instrumented end to end.

I would welcome a discussion about where Thermo Fisher currently feels the greatest architecture friction: duplicated components, review delay, unclear decision rights, cost opacity, operational ownership, or the gap between successful pilots and durable adoption.

Sincerely,

Russell Dudek